

Sarsa – on policy temporal difference algo

```

Initialize  $Q(s, a)$  arbitrarily
Repeat (for each episode):
  Initialize  $s$ 
  Choose  $a$  from  $s$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
  Repeat (for each step of episode):
    Take action  $a$ , observe  $r, s'$ 
    Choose  $a'$  from  $s'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$ 
     $s \leftarrow s'; a \leftarrow a'$ 
  until  $s$  is terminal
  
```

Action for update
as used for trajectory

Figure 6.9 Sarsa: An on-policy TD control algorithm

$$\Delta Q(s, a) = \eta [r - (Q(s, a) - \gamma Q(s', a'))]$$

From: Reinforcement Learning, Sutton and Barto 1998

$$\Delta Q(s, a) = \eta [r - (Q(s, a) - \gamma Q(s', a'))]$$

SARSA algo

SARSA algo

Initialise Q values


Start from initial state s

- 1) Being in state s
choose action a
according to policy π
- 2) Observe reward r
and next state s'
- 3) Choose action a' in state s'
according to policy π
- 4) Update

$$\Delta Q(s, a) = \eta [r - (Q(s, a) - \gamma Q(s', a'))]$$

- 5) $s' \rightarrow s; a' \rightarrow a$

- 6) Goto 1)



$$\Delta Q(s, a) = \eta [r - (Q(s, a) - \gamma Q(s', a'))]$$